

DS 7331 Data Mining: An Introduction to Data Mining and Related Topics

Course Overview

This class introduces the processes of managing, exploring, visualizing, and acting on large amounts of data. This course provides an introduction to data-mining techniques (classification, regression, association mining, clustering, and user/item recommendation) used in analytics. All material covered will be reinforced through hands-on experience using state-of-the art tools to design and execute data mining processes. Class examples will come from the programming language Python (and some examples from R). Pre-requisite courses for this class include basic statistics and probability, linear algebra, calculus, and introductory algorithm analysis (or desire to learn quickly). Experience with databases is helpful but not required.

Assignments will require the python programming language but can also be completed in R or SAS (mostly). Assignments are due in the format of “reports” that include code, visualizations, and summarize all analysis with text or video. **Please note that these assignments are quite comprehensive—student feedback has revealed that they typically take more preparation time than previous courses in data science you may have taken.** An example report and the graded feedback for the report will also be given prior to the submission of the first assignment. This is by design because the assignments are the major components of the course grade.

Learning Objectives

This course is constructed to help students design and use machine learning and data mining techniques. Students will hone their abilities to analyze data, visualize and explain data, and predict outcomes using various learning algorithms. Various techniques for mining rules and building production ready recommendation systems will be discussed. Tools for working with massively data parallel systems will be discussed, as well as techniques for working with massive data sets out of core. While this course is designed for breadth of topics, some topics will be discussed in depth (such as support vector machines). Finally, students will learn to communicate ideas about these technical areas effectively.

Students will:

- Data analysis in python (scikit-learn, pandas, dato, and jupyter)
- Visualization using matplotlib, seaborn, and plotly
- Feature dimension reduction and feature manipulation
- Introduction to common features used in Data Mining
- Linear and non-parametric regression
- Classification techniques (logistic regression, neural networks, SVMs, nearest neighbors, decision trees, ensemble methods)

- Gradient-based optimization techniques (batch, mini-batch, and stochastic gradients)
- Clustering techniques (k-means, DBSCAN, spectral clustering)
- Recommender systems (collaborative filtering, latent factors, matrix factorization)
- Association rule mining
- Parallelism complications are also presented

Textbooks and Materials

Required Reading Text:

- Introduction To Data Mining by Tan, Steinbach, and Kumar
 - First Edition, ISBN-10: 0321321367
- Please note that I am trying to remove a “required textbook” for this course. Most information I present will come from an array of disparate sources. While many theoretical elements come from this text, it is best used for review of video content and examples.

Optional Reading Text:

- Python Machine Learning, by Sebastian Raschka
 - First Edition, ISBN: 978-1-78355-513-0
 - Link: <http://www.amazon.com/Python-Machine-Learning-Sebastian-Raschka/dp/1783555130>
- The Elements of Statistical Learning: Data Mining, Inference, and Prediction, by Trevor Hastie, Robert Tibshirani, and Jerome Friedman
 - Second Edition, ISBN: 978-0-387-84857-0, e-ISBN: 978-0-387-84858-7
 - Protected Link to eBook: https://www.dropbox.com/s/5ychxomsgl085as/statistical_learning_book.pdf?dl=0
- Applied Predictive Modeling by Max Kuhn and Kjell Johnson
 - ISBN: 978-1-4614-6848-6, e-ISBN: 978-1-4614-6849-3
 - Protected Link to eBook: <https://www.dropbox.com/s/oe34ijzti7jg9l/springer%20book.pdf?dl=0>
- Mining of Massive Datasets by Jure Leskovec, Anand Rajaraman, and Jeffrey D. Ullman
 - Link to Free text: [https://www.dropbox.com/s/owd3jff6t0jfaee/MassiveDatasets.pdf? dl=0](https://www.dropbox.com/s/owd3jff6t0jfaee/MassiveDatasets.pdf?dl=0)

Course Github Page: <https://github.com/eclarson/DataMiningNotebooks>

Grading

Assignment/Assessment	Weight on Final Grade
Lab Assignments (3 labs @ 20% each, 1 mini-lab @ 10%)	70%
Live Session Assignments (5 @ 6% each)	30%

Grading Scale

Letter Grade	Percentages
A	93.00 – 100
A-	90.00 – 92.99
B+	87.00 – 89.99
B	83.00 – 86.99
B-	80.00 – 82.99
C+	77.00 – 79.99
C	73.00 – 76.99
C-	70.00 – 72.99
D+	67.00 – 69.99
D	63.00 – 66.99
D-	60.00 – 62.99
F	Else

Team coding is highly encouraged in this course. **All assignments can be completed as a team.** For team portions of assignments, one grade will be issued to the entire team, not individually (that is, the team will sink or swim together for that portion of the assignment). This helps to prevent any issues cropping up about copying/cheating from individuals in the class. Teams can change between assignments. Please handle team conflicts within your team in an appropriate manner. For example, in the case where an individual is not contributing to the team effort, that member can be told that he/she will no longer be part of a given team (well in advance of the assignment). This should be used only where appropriate and in a manner that is cohesive with team inclusiveness. Teams will be evaluated based upon their lab assignments and live- session assignments, as follows:

Lab assignments:

- 70% of grade (3 labs @ 20% each, 1 mini-lab @ 10%)
 - Lab One: Visualization and Data Preprocessing
 - Mini-Lab: Logistic Regression and SVMs
 - Lab Two: Classification
 - Lab Three: Clustering, Association Rules, or Recommenders

Live-Session Assignments:

- 30% of grade (5 at 6% each)
 - Linear Regression in Python
 - Decision Trees and Splitting Criteria
 - KNN, Naive Bayes, and Document Classification
 - Clustering Spatial Data
 - Association Rules in the Titanic

Lab Assignments Grading Schema: In each lab assignment, specific deliverables are asked and should be completed to the best of your ability. Each deliverable will be worth a certain percentage of the grade and you will be graded in terms of the quality of your analysis. Comment code so that it is readable and, where possible, immediately understandable. The

sum total of the these deliverables will be 90% of the points possible for each assignment. If you complete all the project deliverables satisfactorily you should expect a grade of 90%.

The remaining 10% of the points are reserved for **exceptional work that is above and beyond in one or more elements of the analysis**. The instructor will provide some ideas regarding exceptional work for each lab assignment, but students have free reign to provide whatever additional analysis they prefer. Grading of the exceptional work does not follow a strict rubric. It is subjective to the instructor and at the instructor's discretion.

Live-Session Assignment Grading Schema: These assignments are completed under time pressure. I understand this and will grade them accordingly. It is not my wish for these assignments to be “evaluations” of your learning, but, instead, be tools for motivating you to understand the material. **Use these assignments to ask questions to the instructor and get one-on-one feedback!** Because of time pressure, I will drop the lowest live-session assignment grade and replace it with the average of the remaining four. This allowance will not be given for live-session assignments that are not completed (for example, as the result of an unexcused absence). In other words, a 0% cannot be replaced or dropped.

Resubmissions: *Live-session assignments cannot be resubmitted under any circumstances.* However, this is a challenging class taken at an advanced pace. Because of this, lab assignments can be resubmitted if you are unhappy with your performance. Resubmissions have some strict stipulations:

1. Resubmissions are due one week after receiving your grade and instructor feedback.
2. In a resubmission, you should fix the area(s) that the instructor pointed out were deficient. For each deliverable you can recover up to 90% of the points for that deliverable in a resubmission. Do not change deliverables that you already have received 90% or more of the points for as you cannot receive any more credit for these deliverables.
3. You should visibly highlight parts of the submission that are new and summarize the changes at the top of the resubmitted document. Changes to the document that are not documented at the top will not be regraded.
4. Resubmissions are not eligible for “exceptional work.” Any points given in the original assignment for exceptional work will be removed from resubmitted assignment. This is done to disincentive students from performing exceptional work at the expense of the deliverables in the assignment.
5. Because of stipulations 2 and 4 the typical maximum grade for resubmissions is 80%. However, it is possible to get above 80%. This happens when you have received full credit for some deliverables in your original submission and you do not change them.
6. The final lab assignment cannot be regraded because there is no time to resubmit/regrade in the final week of the course.
7. The instructor **reserves the right to not allow resubmission** for assignments that are so poor or so convoluted that it appears a “reasonable effort” was not made. Labs turned in late are also not eligible for resubmission.

Because of the strict stipulations in resubmitting, it is often not advisable to resubmit assignments. However, I believe it can be a very valuable learning experience to try at something you failed at originally. It is my belief that not allowing resubmission after feedback is given tells you “this assignment wasn’t that important anyway.” The stipulations above are meant to guard against students trying to “game” the resubmit system (not penalize

resubmitting). Even still, the best way to get a good lab grade is to effectively complete the assignment in the first submission.

Curves: A dividing line for converting percentages to letter grades must be taken in order to be in line with an academic institution's grading schema. The dividing lines for this course are listed below. Even so, please do not measure your performance in this class by the letter grade you receive. Your potential employers might care very little that you got a B or C in this class if you can knowledgeably talk about data mining systems. Also, I am happy to write recommendation letters for those that performed well in the class (regardless of your grade—performance is different than a letter grade).

Assignment and Assessment Information

Lab Assignments: Periodically lab assignments will be submitted electronically. Lab assignments can be completed in teams of two or three. Late labs will not be accepted and are not eligible for resubmission (see next section). Three lab assignments are due throughout the course and one “mini-lab assignment” worth fewer points (see next section for grading). Lab assignments should be turned in as **one of the three** following options:

1. Rendered Jupyter notebooks (*.ipynb files), **Most preferred method**
2. Embedded HTML (*.htm or *.html) rendering of the Jupyter notebooks (with all images either in a zipped directory with the master HTML file or embedded in the HTML). Please use “zip” to compress directories (i.e., try to avoid *.tar or *.rar because they are harder to view in the online system)
3. PDF documents in report format with all visualizations, code, and analysis included

Lab assignments are posted at the beginning of the course. Most assignments are turned in during a week where formal video/lecture does not take place (except for the mini-lab project). Use this extra time to complete extended or time consuming analyses of the data. There is a high expectation for these assignments. Comment all code **and explain reasoning in detail** in the document. An example graded assignment will be given before the first lab is due.

Live-Session Assignments: Five times during the course we will hold a “Flipped Lecture Assignment” where videos are posted/watched before the live session and we will use class time to complete a live coding assignment. The specifications for the assignment will be given at the start of session and the assignment will then be turned in at the end of session. Students will work in assigned teams and turn in the live-session assignments at the end of the session. Come prepared to work! Live session assignments will be given in the form of a Jupyter notebook with blank spaces for questions and blank spaces for code to be filled in. The live session assignment is completed in three parts, as follows:

1. An initial notebook with a few questions will be given out. This notebook should be answered before coming to class the day of the assignment.
2. The instructor will work out the problems with you during the first 15 minutes of class.

- The remainder of the assignment will then be handed out and teams can work on the assignment until the end of class.

As a team, you will answer the questions to the best of your ability. Live-session assignments should be turned in as rendered Jupyter Notebooks (exactly like the preferred method for lab assignments). These assignments are meant to be formative and facilitate one-on-one interactions with the instructor. Ask questions often and solicit feedback on given answers from the instructor!

If working as a team, all team members will need to be present to receive a grade. An absence during a “live-session assignment” cannot be made up after the assignment has been given in session. However, certain arrangements (in restricted circumstances) can be made beforehand to take the assignment early. For example, known cleared absences like conference/work-related travel, travel related to university sanctioned extracurricular activities, etc. Exceptions will be made for unexpected circumstances such as death/hospitalization of a family member.

Weekly Schedule

Unit	Topic	Readings	Assignments due
1	Introduction to Data Mining in Python	<p>Required Reading/Installation</p> <p>Chapter I and Chapter II from Tan, Steinbach, and Kumar</p> <p>Python Installation: https://www.python.org/</p> <ul style="list-style-type: none"> - You can use with Python 2 or Python 3 for this course. - Install Jupyter notebook on your work station: Jupyter Notebook (interactive python) - Install Pandas <p>Recommended Reading</p> <p>Pandas ecosystem introduction:</p> <ul style="list-style-type: none"> - 10 minute tour: http://vimeo.com/59324550 - http://pandas.pydata.org/pandas-docs/version/0.15.2/tutorials.html 	<p>Install all required packages on your development system (Linux, Mac or, if need be, Windows)</p> <p>Have an idea about the dataset you would like to use during the semester</p> <p>Dataset selection: have a brief explanation of your data, how you plan to preprocess it, and why it is useful for third parties</p>
2	Data Exploration in Python	<p>Required Reading</p> <p>Appendix B and Chapter III from Tan, Steinbach, and Kumar</p> <p>Look at the initial steps in the CRISP-DM framework: http://en.wikipedia.org/wiki/Cross-Industry_Standard_Process_for_Data_Mining</p> <p>Install scipy, numpy, scikit-learn, and matplotlib (these are included with anaconda)</p>	<p>Make sure you have formed teams for the group lab assignments (groups can be size up to three)</p> <p>Start analyzing your dataset using simple visualizations and</p>

		<p>Recommended Reading</p> <p>Data Science Blog by Sebastian Raschka</p> <p>Other visualization software (built to integrate with matplotlib in python):</p> <p>Seaborn</p> <p>MPLD3</p> <p>Plotly (now open source!!)</p> <p>Running aggregated statistics (approximating the median) out-of-core: http://www.silota.com/site-search-blog/approximate-median-computation-big-data/</p> <p>Other python visualization examples:</p> <p>http://pandas.pydata.org/pandas-docs/stable/visualization.html</p> <p>http://matplotlib.org/examples/index.html</p> <p>http://nbviewer.ipython.org/github/mwaskom/seaborn/blob/master/examples/plotting_distributions.ipynb</p>	<p>statistics using Pandas and a visualization package like matplotlib</p>
3	Project Work Week	Turn in the first project work week assignment by Sunday at midnight following the live session	Project 1 due Sunday Following Class
4	From Linear Regression to Classification	<p>Required Reading</p> <p>Appendix D: Linear Regression, from Tan, Steinbach, and Kumar</p> <p>Appendix E: Optimization, from Tan, Steinbach, and Kumar</p> <p>Chapter V (5.4, Artificial Neural Networks), from Tan, Steinbach, and Kumar</p> <p>Recommended Reading</p> <p>Andrew Ng's Linear Regression Overview, CS229 Notes 1: Part I, Sections 1-3</p> <p>https://www.dropbox.com/s/eriypxr001pfso6/cs229-notes1.pdf?dl=0</p>	First Live Session Assignment
5	Logistic Regression and Support Vector Machines	<p>Required Reading</p> <p>Chapter V (5.5 SVMs), from Tan, Steinbach, and Kumar</p> <p>Andrew Ng's Linear Regression Overview, CS229 Notes 1: Part II, Sections 5 and 6 only</p> <p>https://www.dropbox.com/s/eriypxr001pfso6/cs229-notes1.pdf?dl=0</p> <p>Recommended Reading</p> <p>Andrew Ng's SVM Overview, CS229 Notes 3, Section 1-8</p> <p>https://www.dropbox.com/s/60pbyjos01mj3zj/cs229-notes3.pdf?dl=0</p> <p>Chapter Four (4.4, Logistic Regression only) of The Elements of Statistical Learning: Data Mining, Inference, and Prediction, Trevor Hastie, Robert Tibshirani, and Jerome Friedman</p>	Mini Project due, SVM and LR

6	Decision Trees	<p>Required Reading Chapter IV (4.1-4.3, Decision Trees), from Tan, Steinbach, and Kumar scikit-learn decision trees documentation, http://scikit-learn.org/stable/documentation.html (tips, algorithms, mathematical formulation)</p> <p>Recommended Reading Chapter 14 (Classification Trees), from M. Kuhn and K. Johnson, Applied Predictive Modeling Notebook 5 on the course github page (decision trees)</p>	Second Live Session Assignment
7	Model Evaluation and Ensemble Classification	<p>Required Reading Chapter IV (4.4-4.6, Overfitting, Evaluation, Comparing Classifiers), from Tan, Steinbach, and Kumar Chapter V (5.6-5.7, Ensemble Methods and Class Imbalance), from Tan, Steinbach, and Kumar</p> <p>Recommended Reading Chapter Seven and Chapter Fifteen of The Elements of Statistical Learning: Data Mining, Inference, and Prediction, Trevor Hastie, Robert Tibshirani, and Jerome Friedman Chapter 11, from M. Kuhn and K. Johnson, Applied Predictive Modeling • Pipelines in sklearn: http://scikit-learn.org/stable/modules/pipeline.html</p>	Midterm
8	Alternative Classifiers	<p>Required Reading Chapter V (5.2 and 5.3, KNN and Naive Bayes), from Tan, Steinbach, and Kumar scikit-learn documentation for KNN using KD-trees and Naive Bayes: http://scikit-learn.org/stable/modules/naive_bayes.html http://scikit-learn.org/stable/modules/neighbors.html</p> <p>Recommended Reading Chapter 13 (KNN only) of The Elements of Statistical Learning: Data Mining, Inference, and Prediction, Trevor Hastie, Robert Tibshirani, and Jerome Friedman Chapter 13 (KNN and Naive Bayes), from M. Kuhn and K. Johnson, Applied Predictive Modeling Gupta, Maya R., Eric K. Garcia, and Erika Chin. "Adaptive local linear regression with application to printer color management." Image Processing, IEEE Transactions on 17, no. 6 (2008): 936-945.</p>	Third Live Session Assignment

		http://www.countbayesie.com/blog/2015/2/18/bayes-theorem-with-lego Look at the grand Poobah Notebook to see how KNN and Naive Bayes can be used in scikitlearn	
9	Project Work Week	None	Project Two Due Sunday Following Class
10	Basic Clustering Methods	Required Reading Chapter VIII, from Tan, Steinbach, and Kumar scikit-learn documentation on clustering: http://scikit-learn.org/stable/modules/clustering.html Only look at k-means, hierarchical agglomerative clustering, DBSCAN, and cluster validity, That is, skip spectral clustering, birch, etc. Recommended Reading Chapter Thirteen (K-means only) of The Elements of Statistical Learning: Data Mining, Inference, and Prediction, Trevor Hastie, Robert Tibshirani, and Jerome Friedman	Fourth Live Session Assignment
11	More Clustering Methods	Required Reading Chapter IX (9.1, 9.2.2, and 9.4.1-9.4.4, Characteristics and Graph based clustering), from Tan, Steinbach, and Kumar scikit-learn's documentation of GMMs: http://scikit-learn.org/stable/modules/mixture.html scikit-learn's documentation of using the BIC: http://scikit-learn.org/stable/auto_examples/mixture/plot_gmm_selection.html#example-mixture-plot-gmm-selection-py Recommended Reading http://www.ics.uci.edu/~smyth/courses/cs274/notes/EMnotes.pdf Chapter Eight (8.2-8.5, Expectation Maximization and Bayes) of The Elements of Statistical Learning: Data Mining, Inference, and Prediction, Trevor Hastie, Robert Tibshirani, and Jerome Friedman	None
12	Association Rule Mining	Required Reading/Installation Chapter VI, from Tan, Steinbach, and Kumar Install R, and libraries arules and arulesviz http://www.r-project.org Here is a good blog on installing packages with R: http://www.r-bloggers.com/installing-r-packages/ Install Rpy2 (windows users may not be able to install this correctly) Recommended Reading http://cran.r-project.org/web/packages/arules/index.html	Fifth Live Session Assignment

		http://cran.r-project.org/web/packages/arules/arules.pdf http://aimotion.blogspot.com/2013/01/machine-learning-and-data-mining.html Using R and the apriori package: https://www.youtube.com/watch?v=DQGJhZNhG4M Here is a Jupyter notebook I made based on an example from Dr. Michael Hahsler. This might be useful for those wanting to do their final project on association rule mining: Notebook 12 on the course github page (association analysis)	
13	Recommender Systems	Required Reading Look at dato's online documentation https://dato.com/products/create/docs/ In particular, look at the scalable data structure (SFrame) and also the recommender API or examples in their notebooks Look at the open source core of Dato: https://dato.com/products/create/open_source.html Recommended Reading/Installation http://guidetodatamining.com/guide/ch2/DataMining-ch2.pdf http://michael.hahsler.net/research/RecommenderSMU2011/slides/Recomm_2011.pdf https://dato.com/learn/gallery/index.html Install Dato's graphlab create on your system: https://dato.com/products/create/	Begin project
14	Data Parallel and Out of Core	Required Reading: None Recommended Reading: Data parallel in Association Rules and Clustering: http://users.eecs.northwestern.edu/~yingliu/papers/para_arm_cluster.pdf SVMs in parallel (approximate): http://infolab.stanford.edu/~echang/CIKM-tutorial.pdf Chapter Two and Twelve of Mining of Massive Datasets by Jure Leskovec, Anand Rajaraman, and Jeffrey D. Ullman The professional version of Graphlab allows configuring parallel environments on things like AWS and hadoop. If the money is there, these are excellent and scalable implementations that can be run in just a few lines of code. But be sure you actually need parallelism. Often its just out-of-core computation that is needed.	Project due
15	Project Work Week	None	Final Project Due Sunday Following Class

Attendance Policy

Class session attendance is required. Students with three absences or less, who actively participate in class, will not receive any deductions for his/her absence. Starting with the fourth absence, 2% points from the final grade will be deducted for each absence (over the initial three absences). Please note: Rarely are these measures needed!

Excused Absences for University Extracurricular Activities: Students participating in an officially sanctioned, scheduled University extracurricular activity should be given the opportunity to make up class assignments or other graded assignments missed as a result of their participation. It is the responsibility of the student to make arrangements with the instructor prior to any missed scheduled examination or other missed assignment for making up the work. (See the University Undergraduate Catalog for details.)

Please note that this syllabus is subject to change. Any changes to the syllabus will be announced.

Religious Observance: Religiously observant students wishing to be absent on holidays that require missing class should notify their professors in writing at the beginning of the semester, and should discuss with them, in advance, acceptable ways of making up any work missed because of the absence. (See University Policy No. 1.9.)

Students with Disabilities

Students needing academic accommodations for a disability must first be registered with Disability Accommodations & Success Strategies (DASS) to verify the disability and to establish eligibility for accommodations. Students may call 214-768-1470 or visit <http://www.smu.edu/alec/dass.asp> to begin the process. Once registered, students should then schedule an appointment with the professor to make appropriate arrangements.

Honor Code and Academic Integrity

Cheating: Cheating of any kind such as plagiarism or direct copying is strictly prohibited and against the SMU honor code. However, collaboration is strongly encouraged. All lab assignments can be done as a group and turned in as a group.

A note on cheating: You can and should look online for other's that have analyzed data in Jupyter notebooks. However, if you start with their code or use it, you should 100% make it clear where you got it and what you changed (some of the best performers in this class have taken this approach). However, the reference to other work should be included **directly next to the code or visualization** that you are using (that is, including a list of references at bottom of the document and not making it clear where you got parts of code/images/etc. will be considered cheating and plagiarism). Don't try to make it such that the cheating is "plausibly deniable." That is not a defense for behaving unethically and could result in expulsion from the SMU program. When quoting directly from another source, please place that text in quotes and cite the quote immediately. Failure to place text in quotes from other sources will be considered cheating.

I do not mess around here. I have seen a lot of questionable copying of code/text without attribution in Computer Science in recent years. Please attribute where it came from! Many times this copying is the result of not allocating enough time to complete the assignment, or a fear that you cannot get a good grade if you let me know that you are struggling with basic concepts. **It is better to try and fail in these circumstances than to turn in unattributed code.** In the event of plagiarism or cheating, I will turn the matter over to SMU's honor code committee, which could result in expulsion from the University. That is a daunting consequence, especially when I allow resubmissions for the course.

Another note on cheating: I understand small mistakes. Like having a small "similar" snippet of code that you forgot to copy the link to in your assignment (especially when your past performance indicates this was just an "oops, I forgot!"). I am not in the business of failing people for minor missteps, nor would I want to be. These measures are in place for egregious or obvious violations. Most of the assignment in this course, because they are contemporary issues, are hard to cheat.

This syllabus is only a guideline and is not a legal contract. The professor of record for the course has final say on any policies, due dates, etc.

Course Overview and Schedule

Please find the schedule for the course below. Each week is marked with the section titles for videos, along with the number of questions. Graded assignments are marked for the week they are due. The key for the outline is as follows:

- V = Video with slides, length is also given
- S = Screen capture, typically of coding examples, length is also given
- LB = Light Board or white board derivations or examples, length is also given
- Y = YouTube video (of myself or guest instructor)
- Q = Multiple Choice or Short Answer Question, Number of questions is also given
- Q-e = Essay style question
- Q-u = Upload question

1 Introduction to Data Mining in Python

- 1.3 Why Mine Data (V8:29 and 7 Qs)
- 1.4 Types of Data (V1:35, V3:22, 7 Qs, V2:28, 4Qs)
- 1.5 Representing Features in Python Pandas (S35:18)
- 1.6 Feature Manipulation (S13:59)

Total: 66 minutes

Live: Introductions and Numpy through Jupyter Notebook

2 Data Exploration in Python

- 2.3 Types of Data Summary Plots (V2:35, V2:08, V2:03, V2:12, V1:57)
- 2.4 Plotting Made Easy (S28:21)
- 2.5 PCA/LDA for Continuous Variables (V11:47, S20:12)

Total: 70 minutes

Live: Example Student Assignment

3 Project Work Week

Live: Individual Team Meetings

Project One Due Sunday Following Class

4 From Linear Regression to Classification

- 4.3 Linear Regression (LB5:23, LB8:08, LB7:52)
- 4.4 Gradient Based Optimization (LB8:37)
- 4.5 A New Objective: Classification (LB6:36, LB4:44)
- 4.6 Logistic Regression Optimization Setup (LB4:39)

Total: 45 minutes

Live: —First Live Session Assignment—

5 Logistic Regression and Support Vector Machines

- 5.3 Optimizing Two Class Logistic Regression (LB18:13)
- 5.4 Multiclass Logistic regression (LB5:14)
- 5.5 Gradient Optimization: Stochastic, Mini-Batch, Batch (LB7:22)
- 5.6 Support Vector Machines (LB15:14, LB14:43)
- 5.7 Kernel Tricks (LB12:26)

Total: 66 minutes

Live: SVM, LR, and SGD Notebook

Mini Project Due, SVM and LR

6 Decision Trees

6.3 What is a Decision Tree (V2:45, 2Qs, V1:40, 1Q)

6.4 Building a Decision Tree (V2:29, 1Q, V2:03, 1Q, V5:55, 1Q)

6.5 Criteria for Determining the Best Split (V3:36, 1Q, V3:04, 1Q, V3:11, V3:13, 2Q, V2:35)

6.6 Decision Boundaries and Stopping Criteria (V2:56, 1Q, V0:59, 1Q, V3:05, 1Q)

6.7 Decision Tree Generalization (V3:14, V6:22, 2Q, V1:43)

Total: 50 minutes

Live: — *Second Live Session Assignment* —

7 Model Evaluation and Ensemble Classification

7.3 Model Evaluation (V2:28, 1Q-e, V2:35, 1Q, V2:34, 2Q-e)

7.4 Methods for Estimating Performance (V3:37, V1:34, 1Q, V2:10, 1Q, V1:29, 1Q, V4:53)

7.5 Model Comparison (LB2:48, LB5:21, LB4:38, LB2:18)

7.6 Ensemble Classifiers (V4:22, 1Q, V7:52, 1Q, V5:48, 2Q)

Total: 55 minutes

Live: Grand Poobah Classification Notebook

8 Alternative Classifiers

8.3 Nearest Neighbors Classification (V5:21, 1Q, V3:19, 2Q-e, Y21:44, V3:44)

8.4 Nearest Neighbor Regression (V4:58, 1Q)

8.5 KD Trees (V2:36, 1Q, V2:51, 1Q, V5:18)

8.6 Naive Bayes (V1:46, 1Q, V4:53, 1Q, V4:21, 1Q, V1:45)

Total: 60 minutes

Live: — *Third Live Session Assignment* —

9 Project Work Week

Live: Individual Team Meetings

Project Two Due Sunday Following Class

10 Basic Clustering Methods

10.3 What is clustering (V2:18, 2Q, V4:37, 2Q)

10.4 K-means Clustering (V1:13, 1Q, V5:07, 1Q-u, V1:58, V4:27, 2Q, V3:22, V2:32, 1Q)

10.5 Hierarchical Agglomerative Clustering (HAC) (V6:40, 1Q, V2:26)

10.6 Density Based Clustering (V2:30, 1Q, V4:36)

10.7 Clustering Examples (S23:47)

10.8 Clustering Validity (V4:46, V5:06)

Total: 50 minutes

Live: — *Fourth Live Session Assignment* —

11 More Clustering Methods

11.3 Gaussian Mixtures (V2:01, 2Q, V1:28, 1Q, V3:45, 1Q, V3:43)

11.4 Graph Clustering (V3:37, 1Q-u, V2:23, V2:12, V4:14)

11.5 Spectral Clustering (V3:15, V4:53)

Total: 30 minutes

Live: Spectral and GMM Clustering Notebook

12 Association Rule Mining

12.3 What is Association Rule Mining (V4:11, 1Q, V1:45, 1Q, V1:01)

12.4 Frequent Itemset Generation (V2:02, V0:52, V1:27, V1:50, 1Q, V1:29, V3:05, V2:51, V1:35)

12.5 Frequent Pattern Trees (V3:46, 1Q, S15:31, 1Q, S5:11)

12.6 Rules Generation (V3:52, 1Q, V1:35)

12.7 Maximal and Closed Sets (V4:11)

12.8 Evaluating Association Rules (V4:40, 1Q-e, V3:37, V2:42)

Total: 66 minutes

Live: —*Fifth Live Session Assignment*—

13 Recommender Systems (Optional, *Can be over Thanksgiving*)

13.3 Intro to collaborative filtering (V1:33, V3:11, V2:04, V2:41, V1:21)

13.4 SFrames in Dato (Y32:45)

13.5 Recommenders (Y39:32)

13.6 More recommenders in Dato (Opt - Y55:43)

Total: 80 minutes + 55 minutes optional

Live: Recommender Systems Notebook

14 Data Parallel and Out of Core (Optional)

14.3 Data Parallel and Grid Searching

Total: Two Hours, all optional

Live: Typically no Live Session

15 Project Work Week

Live: Individual Team Meetings

Final Project Due Sunday Following Class

Unit One: An Introduction to Data Mining in Python

Lecture Units:

- 1.3 Why Mine Data (V8:29 and 7 Qs)
- 1.4 Types of Data (V1:35, V3:22, 7 Qs, V2:28, 4Qs)
- 1.5 Representing Features in Python Pandas (S35:18)
- 1.6 Feature Manipulation (S13:59)

Total Video: 66 Minutes

Jupyter notebook(s):

- Look at notebook 1 on the course github page.
- Look at notebook 2 on the course github page. (up to visualization)

Unit Objectives

- Understand examples of data mining data sets
- Introduction to attribute types
- Familiarize and install the iPython/Jupyter ecosystem
- Working with data in Pandas: Feature normalization, One hot encoding, Imputation

Required Reading/Installation

- Chapter I and Chapter II from Tan, Steinbach, and Kumar
- Python Installation: <https://www.python.org/>
 - **You can use with Python 2 or Python 3 for this course.**
 - Great package with most of the batteries included: [Anaconda](#)
 - When using anaconda, be sure to use the included “conda” or “pip” package manager
- Install Jupyter notebook on your work station: Jupyter Notebook (interactive python)
 - <http://jupyter.org>
 - this is included in the anaconda distribution
- Install Pandas:
 - [Pydata Pandas Eco-system](#)
 - this is included in the anaconda distribution

Recommended Reading

- Pandas ecosystem introduction:
 - 10 minute tour: <http://vimeo.com/59324550>
 - <http://pandas.pydata.org/pandas-docs/version/0.15.2/tutorials.html>

Unit Assignments

- Install all required packages on your development system (Linux, Mac or, if need be, Windows)
 - Linux and mac will be slightly better operating systems when working with the packages in this class. Although Windows will be fine for the majority of cases, some packages run more slowly or are not optimized for this OS.
- **Have an idea about the dataset you would like to use during the semester**
- During Live Session:
 - Introductions
 - Introduce course grading schema

- Dataset selection: have a brief explanation of your data, how you plan to preprocess it, and why it is useful for third parties
 - Pay attention and start to feel around for others to work with on teams during semester. Teams are due by next week!!
- Come with questions about the Pandas ecosystem and variable representations!
- In class example: what is Numpy and why do we use it? This is example E00 on the course github page.

Dataset Selection for Semester

You will select (as a team) which dataset you wish to work on during the semester. You have a lot of free reign for the analysis, but the dataset must be sufficiently complex. This means:

1. there must be 10 or more attributes for analyzing the data (both continuous and discrete),
2. there must be at least 30,000 records to classify in the data
3. there must be a good documentation of the dataset attributes.

Select a classification dataset from one of the following sources:

- **Public Record.** Chosen from a public repository. Some examples include the Dallas Crime dataset, UCI machine learning repository, and many others.
- **Kaggle.** This website hosts machine learning competitions with prizes to the best performers. You can choose one of the open and enterable competitions-as long as it meets the criteria for complexity outlined above.
 - **Note:** you need to sign up to use Kaggle. Many of the competitions require more than just machine learning to solve. For instance, it is usually up to you to get the dataset in a format that is readable for python. Additionally, you may need to extract features to use (e.g., they give you images and expect you to process the images and extract features).
 - This option is probably more work than the others, but has more reward also.
 - As students, you usually qualify for getting additional prizes and/or job offers, so keep this kind of thing in mind for the future, even if you do not use it for class.
- **Your Own.** If you want to use data from your own research or data that you have collected, that's fine. It still must meet the above mentioned criteria. Verify your dataset with the instructor before choosing this option!

Unit Two: Data Exploration in Python

Lecture Units:

2.3 Types of Data Summary Plots (V2:35, V2:08, V2:03, V2:12, V1:57)

2.4 Plotting Made Easy (S28:21)

2.5 PCA/LDA for Continuous Variables (V11:47, S20:12)

Total Video: 70 minutes

Jupyter notebook(s):

- Look at notebook 2 on the course github page (visualization).
- Look at notebook 3 on the course github page (dimensionality reduction).

Unit Objectives

- Understand how to run aggregated statistics in python Pandas
- Understand how to create and use visualizations with matplotlib and Pandas
 - Exposure to other plotting interfaces like seaborn, MPLD3, and Plotly
- Employ methods for dimensionality reduction in scikit-learn (PCA, LDA, and randomized PCA)

Required Reading/Installation

- Appendix B and Chapter III from Tan, Steinbach, and Kumar
- Look at the initial steps in the CRISP-DM framework: http://en.wikipedia.org/wiki/Cross_Industry_Standard_Process_for_Data_Mining
- Install scipy, numpy, scikit-learn, and matplotlib (these are included with anaconda)

Recommended Reading

- Data Science Blog by [Sebastian Raschka](#)
 - Great intuitions for [PCA](#) and [LDA](#) in python
- Other visualization software (built to integrate with matplotlib in python):
 - Seaborn
 - MPLD3
 - Plotly (now open source!!)
- Running aggregated statistics (approximating the median) out-of-core:
 - <http://www.silota.com/site-search-blog/approximate-median-computation-big-data/>
- Other python visualization examples:
 - <http://pandas.pydata.org/pandas-docs/stable/visualization.html>
 - <http://matplotlib.org/examples/index.html>
 - http://nbviewer.ipython.org/github/mwaskom/seaborn/blob/master/examples/plotting_distributions.ipynb

Unit Assignments

- Make sure you have formed teams for the group lab assignments (groups can be size up to three)
- Look at first project work week assignment (it is due at the end of next week)
 - Start analyzing your dataset using simple visualizations and statistics using Pandas and a visualization package like matplotlib
- During live session:
 - Go over a good example from the past of the first assignment

- Come with questions regarding the visualization and dimensionality notebooks

Unit Three: First Team Project Work Week, Data Exploration and Analysis

- No lecture content: Use Jupyter and pandas ecosystem to explore chosen data sets
 - Turn in the first project work week assignment by Sunday at midnight following the live session
 - During Live Session (optional):
 - Come with questions about using pandas on your datasets
 - I will answer questions to each team individually in breakout sessions, please work on refining assignments as a team during class time
 - Show some of what your team has been working on and any problems you have run into. Others will benefit greatly from these types of questions.
-

Due Sunday before midnight following the Unit 3 live session.

First Project Work Week Assignment



You are to perform analysis of a data set: exploring the statistical summaries of the features, visualizing the attributes, and making conclusions from the visualizations and analysis. Follow the CRISP-DM framework in your analysis (you are not performing all of the CRISP-DM outline, only the portions relevant to understanding and visualization). This report is worth 20% of the final grade. Please upload a report (one per team) with all code used, visualizations, and text in a single document. The format of the document can be PDF, *.ipynb, or HTML. You can write the report in whatever format you like, but it is easiest to turn in the rendered Jupyter notebook.

A note on grading: This lab is mostly about visualizing and understanding your dataset. The largest share of the points is from how you interpret the visuals that you make. Making the visuals is not enough to satisfy each of the rubrics below—you should appropriately explain what the implications of the visualizations are. In other words, expect about 20% of the available points for visuals that have no substantive discussion.

Grading Rubric

- Business Understanding (**10 points total**).
 - Describe the purpose of the data set you selected (i.e., why was this data collected in the first place?). Describe how you would define and measure the outcomes from the dataset. That is, why is this data important and how do you know if you have mined useful knowledge from the dataset? How would you measure the effectiveness of a good prediction algorithm? Be specific.
- Data Understanding (**80 points total**)
 - **[10 points]** Describe the meaning and type of data (scale, values, etc.) for each attribute in the data file.
 - **[15 points]** Verify data quality: Explain any missing values, duplicate data, and outliers. Are those mistakes? How do you deal with these problems? Give justifications for your methods.
 - **[10 points]** Visualize appropriate statistics (e.g., range, mode, mean, median, variance, counts) for a subset of attributes. Describe anything meaningful you found from this or if you found something potentially interesting. Note: You can also use data from other sources for comparison. Explain why the statistics run are meaningful.
 - **[15 points]** Visualize the most interesting attributes (at least 5 attributes, your opinion on what is interesting). **Important: Interpret the implications for each visualization.** Explain for each attribute why the chosen visualization is appropriate.

- **[15 points]** Visualize relationships between attributes: Look at the attributes via scatter plots, correlation, cross-tabulation, group-wise averages, etc. as appropriate. Explain any interesting relationships.
- **[10 points]** Identify and explain interesting relationships between features and the class you are trying to predict (i.e., relationships with variables and the target classification).
- **[5 points]** Are there other features that could be added to the data or created from existing features? Which ones?
- **Exceptional Work (10 points total)**
 - You have free reign to provide additional analyses.
 - One idea: implement dimensionality reduction, then visualize and interpret the results.

Unit Four: From Linear Regression to Classification

Lecture Units:

- 4.3 Linear Regression (LB5:23, LB8:08, LB7:52)
- 4.4 Gradient Based Optimization (LB8:37)
- 4.5 A New Objective: Classification (LB6:36, LB4:44)
- 4.6 Logistic Regression Optimization Setup (LB4:39)

Total Video: 45.5 minutes

A Note to those feeling overwhelmed by the mathematics:

First, your book has good (but compact) explanations of regression and optimization techniques in appendices D and E, respectively. They cover more topics than I do in a short number of pages. Its a great section for reviewing or referencing. We will cover a lot of Appendix E when we hit logistic regression (logistic regression is NOT covered in your book). At the end of the day, we are covering these mathematics because they are extremely useful for understanding how tweaking the different parameters of the algorithms changes generalization and the speed of training. Understanding the ramifications of these mathematics is what might separate a “Good” data scientist from an “Exceptional” data scientist.

There are also great examples of multivariate linear regression on Dr. Andrew Ng's iTunesU page (a Stanford course). <https://itunes.apple.com/WebObjects/MZStore.woa/wa/viewPodcast?id=384233048#ls=1> He uses some of the same notation that I use. I have posted some of Dr. Ng's PDF's for reference, see below.

For logistic regression, you can also see some methods on the Stanford course, but I don't prefer their method of introduction. Instead, I use notation and intuition from Dr. Carlos Guestrin's course on Big Data (from Carnegie Mellon and University of Washington). I will be going over a modified version of Carlos's notation. We will focus only on gradient based optimization for logistic regression. Next week we will hit SVM's and more logistic regression. Your book covers SVM's fairly well, but I will make a few modifications so that their use is more intuitive.

Unit Objectives

- Review the derivation of the linear regression solution
- Understand the ramifications of optimizing for linear regression with a “large number of instances” versus “a large number of attributes”
- Understand how perceptron neural networks build upon the iterative linear regression solution
 - Derive perceptron back-propagation (indirectly)
- Setup derivation of Logistic Regression and Support Vector Machines

Required Reading

- Appendix D: Linear Regression, from Tan, Steinbach, and Kumar
- Appendix E: Optimization, from Tan, Steinbach, and Kumar
- Chapter V (5.4, Artificial Neural Networks), from Tan, Steinbach, and Kumar

Recommended Reading

- Andrew Ng's Linear Regression Overview, CS229 Notes 1: Part I, Sections 1-3
 - <https://www.dropbox.com/s/eriypxr001pfso6/cs229-notes1.pdf?dl=0>

Unit Assignments

- During Live Session
 - **First live session assignment:** Linear Regression
 - You will complete an interactive Jupyter notebook assignment (working in teams if you prefer) during the live session
 - You will work on a Jupyter notebook containing questions about linear regression and linear classification from the videos.
 - An early version of the notebook will be sent out during the week for you to answer the questions before class starts.
 - The instructor will then work out these preliminary questions during the first 15 minutes of class.
 - After this introduction, another notebook will be made available for you to work on as a team. You will be asked to live code (mostly using Numpy) some of the equations we went over in the videos.
 - The live session assignment is meant to reinforce the mathematics with programming so that the concepts become more concrete.
 - You will complete the coding assignment and the questions and upload the updated python notebook at the end of live session.
 - **Ask questions to the instructor! This is meant to facilitate one-on-one time!!**

Unit Five: Logistic Regression and Support Vector Machines

Lecture Units:

- 5.3 Optimizing Two Class Logistic Regression (LB18:13)
- 5.4 Multiclass Logistic regression (LB5:14)
- 5.5 Gradient Optimization: Stochastic, Mini-Batch, Batch (LB7:22)
- 5.6 Support Vector Machines (LB15:14, LB14:43)
- 5.7 Kernel Tricks (LB12:26)

Total Video: 73 minutes

Unit Objectives

- Derive optimization procedure for logistic regression and support vector machines
- Understand alternative optimization methods using iterative gradients: stochastic, batch, and mini-batch gradient descent
- Understand how kernels can be used to represent dot products in infinite dimensional space

Required Reading

- Chapter V (5.5 SVMs), from Tan, Steinbach, and Kumar
- Andrew Ng's Linear Regression Overview, CS229 Notes 1: Part II, Sections 5 and 6 only
 - <https://www.dropbox.com/s/eriypxr001pfso6/cs229-notes1.pdf?dl=0>

Recommended Reading


- Andrew Ng's SVM Overview, CS229 Notes 3, Section 1-8
 - <https://www.dropbox.com/s/60pbyjos01mj3zj/cs229-notes3.pdf?dl=0>
- Chapter Four (4.4, Logistic Regression only) of *The Elements of Statistical Learning: Data Mining, Inference, and Prediction*, Trevor Hastie, Robert Tibshirani, and Jerome Friedman

Unit Assignments

- Finish the introductory assignment for using SVMs and Logistic regression with scikit-learn. It is due Sunday before midnight following live session.
- During live session:
 - Live coding of SVMs and LR in scikit-learn with batch and stochastic gradient descent. The example notebook will be notebook 4 on the course github page.
 - Review how to interpret the weights of logistic regression and linear support vector machines

Due Sunday at Midnight following unit five.

Mini-Project: SVM&LR Classification



You are to perform predictive analysis (classification) upon a data set: model the dataset using methods we have discussed in class: logistic regression and support vector machines, and making conclusions from the analysis. Follow the CRISP-DM framework in your analysis (you are not performing all of the CRISP-DM outline, only the portions relevant to the grading rubric outlined below). This report is worth 10% of the final grade. You may complete this assignment in teams of as many as three people.

Write a report covering all the steps of the project. The format of the document can be PDF, *.ipynb, or HTML. You can write the report in whatever format you like, but it is easiest to turn in the rendered Jupyter notebook. The results should be reproducible using your report. Please **carefully describe every assumption and every step** in your report.

A note on grading: A common mistake I see in this lab is not investigating different input parameters for each model. Try a number of parameter combinations and discuss how the model changed.

SVM and Logistic Regression Modeling

- **[50 points]** Create a logistic regression model and a support vector machine model for the classification task involved with your dataset. Assess how well each model performs (use 80/20 training/testing split for your data). **Adjust parameters of the models to make them more accurate.** If your dataset size requires the use of stochastic gradient descent, then linear kernel only is fine to use. That is, the SGDClassifier is fine to use for optimizing logistic regression and linear support vector machines. For many problems, SGD will be required in order to train the SVM model in a reasonable timeframe.
- **[10 points]** Discuss the advantages of each model for each classification task. Does one type of model offer superior performance over another in terms of prediction accuracy? In terms of training time or efficiency? Explain in detail.
- **[30 points]** Use the weights from logistic regression to interpret the importance of different features for the classification task. **Explain your interpretation in detail.** Why do you think some variables are more important?
- **[10 points]** Look at the chosen support vectors for the classification task. Do these provide any insight into the data? Explain. If you used stochastic gradient descent (and therefore did not explicitly solve for support vectors), try subsampling your data to train the SVC model—then analyze the support vectors from the subsampled dataset.

Unit Six: Decision Trees

Lecture Units:

- 6.3 What is a Decision Tree (V2:45, 2Qs, V1:40, 1Q)
- 6.4 Building a Decision Tree (V2:29, 1Q, V2:03, 1Q, V5:55, 1Q)
- 6.5 Criteria for Determining the Best Split (V3:36, 1Q, V3:04, 1Q, V3:11, V3:13, 2Q, V2:35)
- 6.6 Decision Boundaries and Stopping Criteria (V2:56, 1Q, V0:59, 1Q, V3:05, 1Q)
- 6.7 Decision Tree Generalization (V3:14, V6:22, 2Q, V1:43)

Total Video: 50 minutes

Unit Objectives

- Understand the overall process of Hunt's Algorithm
- Understand the splitting criteria and different methods for splitting on binary, discrete, and continuous attributes
- Get an intuition for how decision trees divide a space (*i.e.*, the decision boundaries)
- Understand the basic overview of different decision tree algorithms: ID3, CART, C4.5/J48
- Understand pruning methods of decision trees

Required Reading

- Chapter IV (4.1-4.3, Decision Trees), from Tan, Steinbach, and Kumar
- scikit-learn decision trees documentation, <http://scikit-learn.org/stable/documentation.html> (tips, algorithms, mathematical formulation)

Recommended Reading

- Chapter 14 (Classification Trees), from M. Kuhn and K. Johnson, *Applied Predictive Modeling*
- Notebook 5 on the course github page (decision trees)

Unit Assignments

- During Live Session:
 - **Second live session assignment:** Decision Trees and Splitting Criteria
 - You will complete an interactive Jupyter notebook assignment (working in teams if you prefer) during the live session
 - You will receive a link to an Jupyter notebook containing questions about decision trees and splitting criteria from the videos.
 - An early version of the notebook will be sent out during the week for you to answer the questions before class starts.
 - The instructor will then work out these preliminary questions during the first 15 minutes of class.
 - After this introduction, another notebook will be made available for you to work on as a team. You will be asked to live code (mostly using Numpy and scikit-learn) some of the equations we went over in the videos.
 - You will complete the coding assignment and the questions and upload the updated python notebook at the end of live session.
 - **Ask questions to the instructor! This is meant to facilitate one-on-one time!**

Unit Seven: Ensembles and Statistical Model Comparison

Lecture Units:

7.3 Model Evaluation (V2:28, 1Q-e, V2:35, 1Q, V2:34, 2Q-e)

7.4 Methods for Estimating Performance (V3:37, V1:34, 1Q, V2:10, 1Q, V1:29, 1Q, V4:53)

7.5 Model Comparison (LB2:48, LB5:21, LB4:38, LB2:18)

7.6 Ensemble Classifiers (V4:22, 1Q, V7:52, 1Q, V5:48, 2Q)

Total Video: 55 minutes

Unit Objectives

- Understand various methods for assessing performance: precision, recall, F-measure, ROC, and AUC
- Understand how to use different cross validations and holdout in scikit-learn
- Statistical model comparisons: confidence intervals in K-fold and holdout
- Understand ensemble classification methods: boosting, bagging, stacking, and cascading
- Use pipelining in scikit-learn for easy chaining of methods

Required Reading

- Chapter IV (4.4-4.6, Overfitting, Evaluation, Comparing Classifiers), from Tan, Steinbach, and Kumar
- Chapter V (5.6-5.7, Ensemble Methods and Class Imbalance), from Tan, Steinbach, and Kumar

Recommended Reading

- Chapter Seven and Chapter Fifteen of *The Elements of Statistical Learning: Data Mining, Inference, and Prediction*, Trevor Hastie, Robert Tibshirani, and Jerome Friedman
- Chapter 11, from M. Kuhn and K. Johnson, *Applied Predictive Modeling*
- Pipelines in sklearn: <http://scikit-learn.org/stable/modules/pipeline.html>

Unit Assignments

- Look at how you might use different classifiers like decision trees and ensemble methods on your chosen data sets
- During live session:
 - Go over the grand Poobah classification notebook of the Olivetti faces dataset
 - Notebook 6 on the course github page (classification)

Unit Eight: Alternative Classification Methods

Lecture Units:

8.3 Nearest Neighbors Classification (V5:21, 1Q, V3:19, 2Q-e, Y21:44, V3:44)

8.4 Nearest Neighbor Regression (V4:58, 1Q)

8.5 KD Trees (V2:36, 1Q, V2:51, 1Q, V5:18)

8.6 Naive Bayes (V1:46, 1Q, V4:53, 1Q, V4:21, 1Q, V1:45)

Total Video: 60 minutes

Unit Objectives

- Understand the complexities associated with k-nearest neighbors classification and regression
 - Employ different distance metrics and understand their advantages
- Understand how data structures can be used to effectively reduce comparisons to the naive “brute force” approach
- Introduce the concept of KNN regression and enclosing/encapsulated k-neighbor regression (e-KNN)
- Understand Naive Bayes classification methods for various types of data

Required Reading

- Chapter V (5.2 and 5.3, KNN and Naive Bayes), from Tan, Steinbach, and Kumar
- scikit-learn documentation for KNN using KD-trees and Naive Bayes:
 - http://scikit-learn.org/stable/modules/naive_bayes.html
 - <http://scikit-learn.org/stable/modules/neighbors.html>

Recommended Reading

- Chapter 13 (KNN only) of *The Elements of Statistical Learning: Data Mining, Inference, and Prediction*, Trevor Hastie, Robert Tibshirani, and Jerome Friedman
- Chapter 13 (KNN and Naive Bayes), from M. Kuhn and K. Johnson, *Applied Predictive Modeling*
- Gupta, Maya R., Eric K. Garcia, and Erika Chin. "Adaptive local linear regression with application to printer color management." *Image Processing*, IEEE Transactions on 17, no. 6 (2008): 936-945.
- <http://www.countbayesie.com/blog/2015/2/18/bayes-theorem-with-lego>
- Look at the grand Poobah Notebook to see how KNN and Naive Bayes can be used in scikit-learn:
 - Included in notebook 6 of the course github page.

Unit Assignments

- During live session:
 - **Third live session assignment:** KNN and Naive Bayes
 - You will complete an interactive Jupyter notebook assignment (working in teams if you prefer) during the live session
 - You will receive a link to a Jupyter notebook that I made containing questions about using KNN and Naive Bayes in scikit-learn.
 - An early version of the notebook will be sent out during the week for you to answer the questions before class starts.

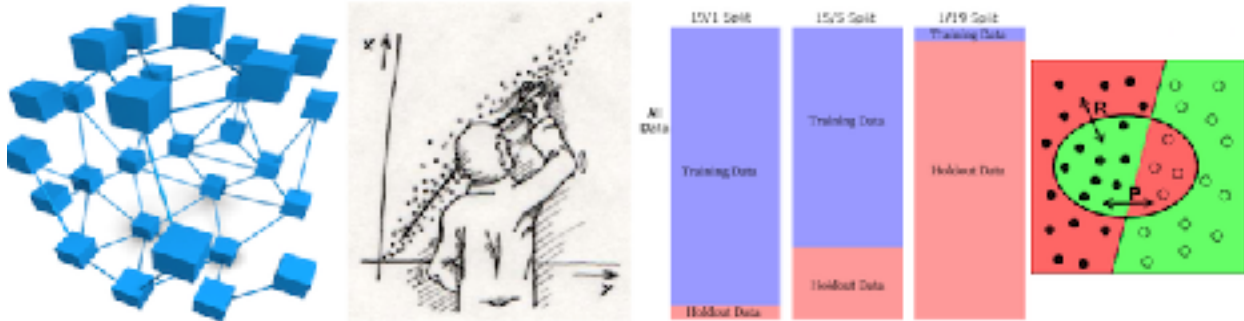
- The instructor will then work out these preliminary questions during the first 15 minutes of class.
- After this introduction, another notebook will be made available for you to work on as a team. You will be asked to live code (mostly using Numpy and scikit-learn) some of the algorithms we went over in the videos.
- You will complete the coding assignment and the questions and upload the updated python notebook at the end of live session.
- **Ask questions to the instructor! This is meant to facilitate one-on-one time!**

Unit Nine: Second Team Project Work Week, Classification

- No lecture content: Use scikit-learn to properly predict from your dataset
 - Turn in the second project work week assignment by Sunday at midnight following live session
 - Please note that this assignment should be much more rigorous than the previous min-project on classification. More explanations and data exploration is required for this assignment.
 - During Live Session (Optional):
 - Come with questions about using pandas on your datasets
 - I will answer questions to each team individually in breakout sessions, please work on refining assignments as a team during class time
 - Show some of what your team has been working on and any problems you have run into. Others will benefit greatly from these types of questions.
-

This is due the Sunday (midnight) following the ninth live session of the course.

Second Project Work Week Assignment



You are to build upon the predictive analysis (classification) that you already completed in the previous mini-project, adding additional modeling from new classification algorithms as well as more explanations that are inline with the CRISP-DM framework. You should use appropriate cross validation for all of your analysis (explain your chosen method of performance validation *in detail*). Try to use as much testing data as possible *in a realistic manner* (you should define what you think is realistic and why).

This report is worth 20% of the final grade. Please upload a report (one per team) with all code used, visualizations, and text in a single document. The format of the document can be PDF, *.ipynb, or HTML. You can write the report in whatever format you like, but it is easiest to turn in the rendered Jupyter notebook. The results should be reproducible using your report. Please carefully describe every assumption and every step in your report.

Dataset Selection

Select a dataset identically to the way you selected for the first project work week and mini-project. You are not required to use the same dataset that you used in the past, but you are encouraged. You must identify two tasks from the dataset to regress or classify. That is:

- two classification tasks OR
- two regression tasks OR
- one classification task and one regression task

For example, if your dataset was from the diabetes data you might try to predict two tasks: (1) classifying if a patient will be readmitted within a 30 day period or not, and (2) regressing what the total number of days a patient will spend in the hospital, given their history and specifics of the encounter like tests administered and previous admittance.

Grading Rubric

- Data Preparation (**15 points total**)
 - **[10 points]** Define and prepare your class variables. Use proper variable representations (int, float, one-hot, etc.). Use pre-processing methods (as needed) for dimensionality reduction, scaling, etc. Remove variables that are not needed/useful for the analysis.

- **[5 points]** Describe the final dataset that is used for classification/regression (include a description of any newly formed variables you created).
- **Modeling and Evaluation (70 points total)**
 - **[10 points]** Choose and explain your evaluation metrics that you will use (i.e., accuracy, precision, recall, F-measure, or any metric we have discussed). Why are the measure(s) appropriate for analyzing the results of your modeling? Give a detailed explanation backing up any assertions.
 - **[10 points]** Choose the method you will use for dividing your data into training and testing splits (i.e., are you using Stratified 10-fold cross validation? Why?). Explain why your chosen method is appropriate or use more than one method as appropriate. For example, if you are using time series data then you should be using continuous training and testing sets across time.
 - **[20 points]** Create three different classification/regression models for each task (e.g., random forest, KNN, and SVM for task one and the same or different algorithms for task two). Two modeling techniques must be new (but the third could be SVM or logistic regression). **Adjust parameters as appropriate to increase generalization performance using your chosen metric.** You must investigate different parameters of the algorithms!
 - **[10 points]** Analyze the results using your chosen method of evaluation. Use visualizations of the results to bolster the analysis. Explain any visuals and analyze why they are interesting to someone that might use this model.
 - **[10 points]** Discuss the advantages of each model for each classification task, if any. If there are not advantages, explain why. Is any model better than another? Is the difference significant with 95% confidence? *Use proper statistical comparison methods.* You must use statistical comparison techniques—be sure they are appropriate for your chosen method of validation as discussed in unit 7 of the course.
 - **[10 points]** Which attributes from your analysis are most important? Use proper methods discussed in class to evaluate the importance of different attributes. Discuss the results and hypothesize about why certain attributes are more important than others for a given classification task.
- **Deployment (5 points total)**
 - How useful is your model for interested parties (i.e., the companies or organizations that might want to use it for prediction)? How would you measure the model's value if it was used by these parties? How would you deploy your model for interested parties? What other data should be collected? How often would the model need to be updated, etc.?
- **Exceptional Work (10 points total)**
 - You have free reign to provide additional analyses.
 - One idea: grid search parameters in a parallelized fashion and visualize the performances across attributes. Which parameters are most significant for making a good model for each classification algorithm?

Unit Ten: Basic Clustering

Lecture Units:

10.3 What is clustering (V2:18, 2Q, V4:37, 2Q)

10.4 K-means Clustering (V1:13, 1Q, V5:07, 1Q-u, V1:58, V4:27, 2Q, V3:22, V2:32, 1Q)

10.5 Hierarchical Agglomerative Clustering (HAC) (V6:40, 1Q, V2:26)

10.6 Density Based Clustering (V2:301Q, V4:36)

10.7 Clustering Examples (S23:47)

10.8 Clustering Validity (V4:46, V5:06)

Total Video: 50 minutes

- Notebook 9 on the course github page (basic clustering)

Unit Objectives

- Be able to identify the various types of clusters that exist (globular, partitional, etc.)
- Understand efficiency of K-means and the different implementations
 - Understand the limitations of k-means
- Understand hierarchical agglomerative clustering and the different merging schemes
- Understand the basics of DBSCAN and how to find Eps given a MinPts
- Understand how to evaluate the validity of clusters using various methods

Required Reading

- Chapter VIII, from Tan, Steinbach, and Kumar
- scikit-learn documentation on clustering:
 - <http://scikit-learn.org/stable/modules/clustering.html>
 - Only look at k-means, hierarchical agglomerative clustering, DBSCAN, and cluster validity. That is, skip spectral clustering, birch, etc.

Recommended Reading

- Chapter Thirteen (K-means only) of *The Elements of Statistical Learning: Data Mining, Inference, and Prediction*, Trevor Hastie, Robert Tibshirani, and Jerome Friedman

Unit Assignments

- During Live Session:
 - **Fourth live session assignment:** Basic Clustering
 - You will complete an interactive Jupyter notebook assignment (working in teams if you prefer) during the live session
 - You will receive a link to a Jupyter notebook that I made containing questions about using k-means, HAC, and DBSCAN in scikit-learn. You will be asked to live code (mostly using scikit-learn) some of the algorithms we went over in the videos.
 - **A note for windows users:** k-means clustering is not optimized in the current release of scikit for the windows OS using python 2. It therefore takes many minutes to run each iteration (instead of milliseconds). The easiest way to run the notebook is via mac or linux or to upgrade to python 3 on windows.
 - You will complete the coding assignment and the questions and upload the updated python notebook at the end of live session.
 - **Ask questions to the instructor! This is meant to facilitate one-on-one time!!**

Unit Eleven: More Clustering

Lecture Units:

11.3 Gaussian Mixtures (V2:01, 2Q, V1:28, 1Q, **V3:45**, 1Q, V3:43)

11.4 Graph Clustering (V3:37, 1Q-u, V2:23, V2:12, V4:14)

11.5 Spectral Clustering (V3:15, V4:53)

Total Video: 30 minutes

Unit Objectives

- Understand limitations of current clustering schemes
- Introduce Expectation Maximization and Probabilistic clustering with mixture models
- Graph based clustering:
 - Understand two common algorithms: spectral clustering and chameleon

Required Reading

- Chapter IX (9.1, 9.2.2, and 9.4.1-9.4.4, Characteristics and Graph based clustering), from Tan, Steinbach, and Kumar
- scikit-learn's documentation of GMMs: <http://scikit-learn.org/stable/modules/mixture.html>
- scikit-learn's documentation of using the BIC: http://scikit-learn.org/stable/auto_examples/mixture/plot_gmm_selection.html#example-mixture-plot-gmm-selection-py

Recommended Reading

- <http://www.ics.uci.edu/~smyth/courses/cs274/notes/EMnotes.pdf>
- Chapter Eight (8.2-8.5, Expectation Maximization and Bayes) of *The Elements of Statistical Learning: Data Mining, Inference, and Prediction*, Trevor Hastie, Robert Tibshirani, and Jerome Friedman

Unit Assignments

- Look for at final assignment for the course and start to make selections about what topic you would like to do: Clustering, Association Rule Mining, or Collaborative Filtering
- During Live Session:
 - Go over the example clustering notebook in class:
 - Notebook 11 on the course github page (advanced clustering)
 - Introduce Dato for clustering and the SFrame from graphlab-create

Unit Twelve: Association Rule Mining

Lecture Units:

- 12.3 What is Association Rule Mining (V4:11, 1Q, V1:45, 1Q, V1:01)
- 12.4 Frequent Itemset Generation (V2:02, V0:52, V1:27, V1:50, 1Q, V1:29, V3:05, V2:51, V1:35)
- 12.5 Frequent Pattern Trees (V3:46, 1Q, S15:31, 1Q, S5:11)
- 12.6 Rules Generation (V3:52, 1Q, V1:35)
- 12.7 Maximal and Closed Sets (V4:11)
- 12.8 Evaluating Association Rules (V4:40, 1Q-e, V3:37, V2:42)

Total Video: 66 minutes

Unit Objectives

- Understand frequent itemset generation and how it can be decoupled from association rule generation
- Understand the apriori algorithm and how it reduces the complexity of finding itemsets that are frequent
- Get an intuition about frequent pattern trees
- Understand rules generation apriori algorithm
- Understand compact rules representations: maximal and closed sets

Required Reading/Installation

- Chapter VI, from Tan, Steinbach, and Kumar
- Install R, and libraries arules and arulesviz
 - <http://www.r-project.org>
 - Here is a good blog on installing packages with R: <http://www.r-bloggers.com/installing-r-packages/>
- Install Rpy2 (windows users may not be able to install this correctly)

Recommended Reading

- <http://cran.r-project.org/web/packages/arules/index.html>
- <http://cran.r-project.org/web/packages/arules/arules.pdf>
- <http://aimotion.blogspot.com/2013/01/machine-learning-and-data-mining.html>
- Using R and the apriori package: <https://www.youtube.com/watch?v=DQGJhZNhG4M>
- Here is a Jupyter notebook I made based on an example from Dr. Michael Hahsler. This might be useful for those wanting to do their final project on association rule mining:
 - Notebook 12 on the course github page (association analysis)

Unit Assignments

- During live session:
 - **Final live session assignment:** Association Rules
 - You will complete an interactive Jupyter notebook assignment (working in teams if you prefer) during the live session
 - You will receive a link to a Jupyter notebook that I made containing questions about using association rule mining in Jupyter using R. You will be asked to live code (mostly

using Numpy and R) for frequent itemsets and rules generation we talked about in the videos.

- **A note for windows users:** Rpy2 ended its support for windows. It therefore will likely not run on the windows OS. The easiest way to run the notebook is via mac or linux. If you cannot install Rpy2, an alternative notebook will be give to you.
- You will complete the coding assignment and the questions and upload the updated python notebook at the end of live session.
- To test if you have everything installed correctly on your system: please download the following notebook with data and try to run it:
 - ICA5_Test folder from the course github page
- **Ask questions to the instructor! This is meant to facilitate one-on-one time!!**

Unit Thirteen: Collaborative Filtering, Optional

Lecture Units:

13.3 Intro to collaborative filtering (V1:33, V3:11, V2:04, V2:41, V1:21)

13.4 SFrames in Dato (Y32:45)

13.5 Recommenders (Y39:32)

13.6 More recommenders in Dato (Opt - Y55:43)

Total Video: 80 minutes + 55 optional

Unit Objectives

- Understand the basics of collaborative filtering using item-item similarity and matrix factorization for user-item similarity
- Understand the utility of out-of-core computations without parallelization
- Understand how to evaluate different recommender systems in graph lab create

Required Reading

- Look at dato's online documentation
 - <https://dato.com/products/create/docs/>
 - In particular, look at the scalable data structure (SFrame) and also the recommender API or examples in their notebooks
- Look at the open source core of Dato:
 - https://dato.com/products/create/open_source.html

Recommended Reading/Installation

- <http://guidetodatamining.com/guide/ch2/DataMining-ch2.pdf>
- http://michael.hahsler.net/research/Recommender_SMU2011/slides/Recomm_2011.pdf
- <https://dato.com/learn/gallery/index.html>
- Install Dato's graphlab create on your system: <https://dato.com/products/create/>

Unit Assignments

- Keep working on your final projects!
- During Live session:
 - Explain evaluation methods for recommender systems
 - Go over collaborative filtering in Dato notebook:
 - Notebook 13 from the course github page (collaborative filtering)

Unit Fourteen: Data Parallel in Python, Optional

Lecture Units:

14.3 Data Parallel (V5:10, V3:08)

14.4 Pycon Tutorial (Y1:33:00)

Total Video: 2 hours

As an alternative to the sklearn video content: Look at Dask, Apple's Turi Create API, or parallelism with PySpark

Unit Objectives

- Understand advantages and limitations of map-reduce in a iterative data mining algorithms
- Understand grid searching and parametrization with scikit-learn

Required Reading

- None

Recommended Reading

- Data parallel in Association Rules and Clustering:
 - http://users.eecs.northwestern.edu/~yingliu/papers/para_arm_cluster.pdf
- SVMs in parallel (approximate): <http://infolab.stanford.edu/~echang/CIKM-tutorial.pdf>
- Chapter Two and Twelve of *Mining of Massive Datasets* by Jure Leskovec, Anand Rajaraman, and Jeffrey D. Ullman
- Apple's Turi Create API allows configuring parallel environments on things like AWS and hadoop. If the money is there, these are excellent and scalable implementations that can be run in just a few lines of code. But *be sure you actually need parallelism*. Often its just out-of-core computation that is needed.

Unit Assignments

- Keep working on the final assignment! It is due at the end of next week!!
 - During Final Live Session (Optional):
 - Come with questions about using clustering, rule mining, or recommendations on your datasets
 - Show some of what your team has been working on and any problems you have run into. Others will benefit greatly from these types of questions.
 - I (and other teams) will answer as many questions as possible during the live session
-

Unit Fifteen: Final Team Project Work Week

- No lecture content: Use scikit-learn/R/Dato to properly cluster/associate/recommend from your dataset
 - Turn in the final project work week assignment by Sunday at midnight following live session
 - Please make this assignment a breadth of what you have learned in the final third of the class, showing me as much as possible in terms of different techniques and proper evaluation on your datasets.
 - During Live Session (Optional):
 - Come with questions about using pandas on your datasets
 - I will answer questions to each team individually in breakout sessions, please work on refining assignments as a team during class time
 - Show some of what your team has been working on and any problems you have run into. Others will benefit greatly from these types of questions.
-

Unit Fifteen: Final Project Work Week



CRISP-DM Capstone: Association Rule Mining, Clustering, or Collaborative Filtering

In the final assignment for this course, you will be using one of three different analysis methods:

- Option A: Use clustering on an unlabeled dataset to provide insight or features
- Option B: Use transaction data for mining associations rules
- Option C: Use collaborative filtering to build a custom recommendation system

Your choice of dataset will largely determine the task that you are trying to achieve. Though the dataset does not need to change from your previous tasks. For example, you might choose to use clustering on your data as a preprocessing step that extracts different features. Then you can use those features to build a classifier and analyze its performance in terms of accuracy (precision, recall) and speed. Alternatively, you might choose a completely different dataset and perform rule mining or build a recommendation system.

Dataset Selection and Toolkits

As before, you need to choose a dataset that is not small. It might be massive in terms of the number of attributes (or transactions), classes (or items, users, etc.) or whatever is appropriate for the task you are performing. Note that scikit-learn can be used for clustering analysis, but not for Association Rule Mining (you should use R) or collaborative filtering (you should use graphlab-create from Dato). Both can be run using Jupyter notebooks as shown in lecture.

- One example of a recommendation dataset is the movie lens rating data: <http://grouplens.org/datasets/movielens/>
- Some examples of association rule mining datasets: <http://fimi.ua.ac.be/data/>

Write a report covering in detail all the steps of the project. The results need to be reproducible using only this report. Describe all assumptions you make and include all code you use in the Jupyter notebook or as supplemental functions. Follow the CRISP-DM framework in your analysis (you are performing all of the CRISP-DM outline). This report is worth 20% of the final grade.

Grading Rubric

Business Understanding (10 points total).

- **[10 points]** Describe the purpose of the data set you selected (i.e., why was this data collected in the first place?). How will you measure the effectiveness of a good

algorithm? Why does your chosen validation method make sense for this specific dataset and the stakeholders needs?

Data Understanding (20 points total)

- **[10 points]** Describe the meaning and type of data (scale, values, etc.) for each attribute in the data file. Verify data quality: Are there missing values? Duplicate data? Outliers? Are those mistakes? How do you deal with these problems?
- **[10 points]** Visualize the any important attributes appropriately. Important: Provide an interpretation for any charts or graphs.

Modeling and Evaluation (50 points total)

Different tasks will require different evaluation methods. Be as thorough as possible when analyzing the data you have chosen and use visualizations of the results to explain the performance and expected outcomes whenever possible. Guide the reader through your analysis with plenty of discussion of the results. Each option is broken down by:

- **[10 Points]** Train and adjust parameters
- **[10 Points]** Evaluate and Compare
- **[10 Points]** Visualize Results
- **[20 Points]** Summarize the Ramifications
- **Option A: Cluster Analysis**
 - *Train:* Perform cluster analysis using several clustering methods (adjust parameters).
 - *Eval:* Use internal and/or external validation measures to describe and compare the clusterings and the clusters— how did you determine a suitable number of clusters for each method?
 - *Visualize:* Use tables/visualization to discuss the found results. Explain each visualization in detail.
 - *Summarize:* Describe your results. What findings are the most interesting and why?
- **Option B: Association Rule Mining**
 - *Train:* Create frequent itemsets and association rules (adjust parameters).
 - *Eval:* Use several measure for evaluating how interesting different rules are.
 - *Visualize:* Use tables/visualization to discuss the found results.
 - *Summarize:* Describe your results. What findings are the most compelling and why?
- **Option C: Collaborative Filtering**
 - *Train:* Create user-item matrices or item-item matrices using collaborative filtering (adjust parameters).
 - *Eval:* Determine performance of the recommendations using different performance measures (explain the ramifications of each measure).
 - *Visualize:* Use tables/visualization to discuss the found results. Explain each visualization in detail.
 - *Summarize:* Describe your results. What findings are the most compelling and why?

Deployment (10 points total)

- Be critical of your performance and tell the reader how you current model might be usable by other parties. Did you achieve your goals? If not, can you reign in the utility of your modeling?
 - How useful is your model for interested parties (i.e., the companies or organizations that might want to use it)?
 - How would you deploy your model for interested parties?
 - What other data should be collected?

- How often would the model need to be updated, etc.?

Exceptional Work (10 points total)

- You have free reign to provide additional analyses or combine analyses.